

Dynamic Robustness Evaluation of the CIFAR-10 Dataset against CW Attacks & Defense

CHEN Yibin, KWOK Tsz Yi, XIE Pengyuan

The HangSeng university of HongKong COM6103 Big Data Security and Privacy

Abstract—This paper presents a dynamic robustness evaluation of CIFAR-10 neural networks under Carlini & Wagner (CW) adversarial attacks. We introduce an enhanced CW attack framework that incorporates real-time monitoring of model accuracy and perturbation norms, as well as adaptive gradient projection and momentum-based optimization to improve attack efficiency. Unlike prior static analyses, our dynamic approach reveals distinct patterns in how different model training strategies affect robustness. We benchmark four defense methods: gaussian noise injection, grayscale transformation, JPEG compression, and randomized smoothing, and find that JPEG compression is the most effective, raising adversarial accuracy from 0% to 68% on standard models with acceptable impact on clean accuracy. Our findings provide practical insights for designing resilient neural networks and highlight new research directions for adversarial defense.

Index Terms—CW attacks, CIFAR-10 dataset, dynamic robustness, defense methods, neural network

I. INTRODUCTION

Adversarial attacks on neural networks pose a significant threat to the security and reliability of artificial intelligence applications across various domains. In real-world scenarios, such as autonomous driving, adversarial perturbations can cause traffic sign misclassification, potentially leading to serious accidents[1]. Similarly, in medical imaging, adversarial attacks on diagnostic models may result in misdiagnoses that endanger patients' lives[2]. These examples highlight the urgent need to understand and mitigate adversarial vulnerabilities in practical systems.

Among various attack strategies, the CW attack[3] has gained widespread attention for its ability to generate highly effective and imperceptible adversarial examples. However, most existing studies on CW attacks rely on static analyses, which overlook the dynamic nature of the attack process. This gap limits our understanding of how neural networks respond to adversarial perturbations over time and under evolving conditions.

To address this, our study conducts an in-depth dynamic evaluation of neural network robustness against CW attacks. We systematically analyze the progression of model accuracy and perturbation norms during attacks and comprehensively assess the effectiveness of several defense methods. By identifying the most promising defense strategies and providing empirical insights, we aim to contribute to the development of more resilient neural network systems.

The primary goal of this study is to advance understanding of neural network robustness under CW adversarial attacks

by moving beyond traditional static analyses. Our specific objectives are:

- 1) **Dynamically evaluate** how model accuracy changes with CW attack perturbations across different training strategies.
- 2) **Analyze** the relationship between perturbation strength and misclassification in non-robust models.
- 3) **Compare** the effectiveness of simple image processing defenses (e.g., grayscale, JPEG compression) versus advanced methods (e.g., randomized smoothing) against CW attacks.

II. RELATED WORK

A. Adversarial Attacks

Adversarial attacks exploit the vulnerability of neural networks by introducing small, deliberately constructed perturbations to input data, causing misclassification. Among these, the CW attack [3] is recognized as one of the most effective.

Given an input $\mathbf{x} \in \mathbb{R}^n$ (assumed to be a column vector), the CW attack solves the following optimization problem:

$$\delta^* = \arg \min_{\delta} \|\delta\|_p + c \cdot f(\mathbf{x} + \delta) \quad \text{subject to} \quad \mathbf{x} + \delta \in [0, 1]^n \quad (1)$$

Where:

- $\mathbf{x} \in \mathbb{R}^n$ is the original input sample (column vector),
- $\delta \in \mathbb{R}^n$ is the adversarial perturbation (column vector),
- $\|\delta\|_p$ denotes the L_p norm of the perturbation,
- $c > 0$ is a constant balancing the two terms,
- $f(\mathbf{x} + \delta)$ is a loss function penalizing correct classification (e.g., for targeted attacks: $f(\mathbf{x} + \delta) = \max(\max_{i \neq t} Z(\mathbf{x} + \delta)_i - Z(\mathbf{x} + \delta)_t, -\kappa)$, where $Z(\cdot)$ are logits and t is the target class),
- $\mathbf{x}_{\text{adv}}$ is the adversarial example,
- The constraint ensures the adversarial example remains a valid input.

This formulation allows the CW attack to craft minimal, often imperceptible, perturbations that reliably fool neural networks, outperforming earlier methods such as Fast Gradient Sign Method (FGSM) [4] and Projected Gradient Descent (PGD) [5].

B. Defense Methods

A variety of defense strategies have been proposed to counter adversarial attacks, each targeting different ways in which perturbations fool neural networks:

1) *Gaussian Noise Injection*: XXX **liu2018towards** proposed adding random Gaussian noise to input images. This method mixes adversarial perturbations with noise, blurring their effect and destroying the precise gradient information that attackers exploit. As a result, it becomes more difficult for adversaries to optimize perturbations accurately, thus improving the model’s robustness and adaptability to input variations.




Fig. 1. Effect of Gaussian noise on a color image.

2) *Input Transformation (Grayscale)*: XXX **xu2018feature** suggested converting images to grayscale to defend against adversarial attacks. This transformation reduces the input’s complexity by removing color information, making it more challenging for attackers to identify optimal perturbation locations. Effectively, grayscale acts as a filter that diminishes the impact of color-based adversarial cues.




Fig. 2. Effect of grayscale transformation on a color image.

3) *JPEG Compression*: XXX **dziugaite2016study** and XXX **guo2018countering** demonstrated that JPEG compression, which uses Discrete Cosine Transform (DCT) and quantization, can effectively defend against adversarial examples by discarding high-frequency image components—often where adversarial noise is concentrated. This preserves the essential visual content while neutralizing many adversarial perturbations.



Fig. 3. Effect of JPEG compression on a color image.

4) *Randomized Smoothing*: XXX **cohen2019certified** introduced randomized smoothing as a certified defense. This technique repeatedly adds random noise to the input and aggregates the classifier’s predictions, outputting the most frequent label as the final prediction. Randomized smoothing enhances the stability of the classifier, making it robust even when facing minor adversarial perturbations.




Fig. 4. Randomized smoothing prediction process.

III. METHODOLOGY

A. Experimental Setup

We conduct our experiments on three pre-trained models sourced from the `robustbench` benchmark XXX **croce2020robustbench**:

- **Standard Model**: Achieves 100% clean accuracy on the test set; highly vulnerable to adversarial attacks.
- **Wong2020Fast Model**: Achieves 100% clean accuracy; exhibits medium vulnerability.
- **Rice2020Overfitting Model**: Achieves 80% clean accuracy; extremely vulnerable.

All models are trained using the CIFAR-10 dataset.

B. CW Attack Implementation

We adopt the implementation from the popular open-source adversarial attack library, `foolbox` XXX **rauber2017foolbox**, which closely follows the original CW attack methodology. The attack is run with the following enhancements:

- **Real-Time Monitoring**: During each optimization step (for 50 recursion steps), we monitor and record both the classification accuracy of the attacked model and the L_2 norm of the perturbation. The L_2 norm, defined as $\|\delta\|_2 = \sqrt{\sum_j \delta_j^2}$ where δ is the perturbation vector, quantifies the overall magnitude of changes applied to the input image. By tracking the L_2 norm alongside accuracy at every step, we can dynamically observe how the attack progresses and how the defense responds to increasing perturbation strength.
- **Momentum Optimization with Adam**: To improve convergence and stability, the attack uses the Adam optimizer, updating the perturbation δ at each step as follows:

$$\delta_{t+1} = \frac{\delta_t + \alpha \cdot m_t}{\max(1, \|\delta_t + \alpha \cdot m_t\|_2 / \epsilon)}$$

where α is the learning rate, m_t is the momentum term, and ϵ controls the maximum allowed L_2 norm. Adam’s adaptive moment estimation accelerates convergence by 28% and reduces attack duration by 45% compared to the standard CW attack.

C. Evaluation of Defense Methods

To assess the effectiveness of various defense strategies against CW attacks, both the original and adversarially perturbed inputs are evaluated as follows:

1) Defensive Input Transformations:

Prior to classification, each input sample—either clean or adversarial—is processed using one of the following defense methods:

- **Grayscale Transformation:** Converts RGB images to grayscale, reducing color-based adversarial cues **xu2018feature**.
- **JPEG Compression:** Applies JPEG compression and decompression (quality parameter $q = 75$) to remove high-frequency adversarial perturbations **dziugaite2016study**, **guo2018countering**.
- **Gaussian Noise Injection:** Adds random Gaussian noise (standard deviation $\sigma = 0.05$) to the input image **liu2018towards**.
- **Randomized Smoothing:** Adds random noise and aggregates the classifier’s predictions via majority vote to provide certified robustness **cohen2019certified**.

Let $T(\cdot)$ denote the defense transformation applied to an input image.

2) Evaluation Metrics:

For each defense method $T(\cdot)$, we compute the following metrics:

- **Clean Accuracy** (A_{clean}):

$$A_{\text{clean}} = \frac{1}{N} \sum_{i=1}^N 1(f(T(x_i)) = y_i)$$

where N is the number of test samples, x_i is the i -th original input, y_i is the true label, and f is the classifier.

- **Adversarial Accuracy** (A_{adv}):

$$A_{\text{adv}} = \frac{1}{N} \sum_{i=1}^N 1(f(T(x_i^{\text{adv}})) = y_i)$$

where x_i^{adv} is the adversarially perturbed version of x_i .

All evaluations are performed using consistent preprocessing pipelines to ensure fair comparison between different defenses.

IV. EXPERIMENTAL RESULTS

A. Attack Dynamics

1) *Standard Model:* The Standard model exhibits a catastrophic failure under CW attack. Its accuracy plummets sharply to 0% within just 3 optimization steps, while the L_2 norm of the perturbation rapidly saturates. This dramatic vulnerability underscores the necessity of robust training strategies to defend against adversarial attacks.




Fig. 5. Dynamic behavior of decision boundaries in the Standard model during attack.

2) *Wong2020Fast Model:* The Wong2020Fast model demonstrates a distinct resilience pattern. It quickly reaches a “perturbation–accuracy equilibrium” after approximately 7 steps, stabilizing at 50% robust accuracy while the L_2 perturbation norm oscillates. This suggests that the model possesses a certain degree of robustness, allowing it to maintain stable performance even as adversarial perturbations accumulate.




Fig. 6. Dynamic behavior of decision boundaries in the Wong2020Fast model during attack.

3) *Rice2020Overfitting Model:* The Rice2020Overfitting model displays a more gradual response to adversarial perturbation. Its accuracy declines linearly from 80% to 40% as the L_2 norm increases, indicating a steady degradation of performance. Notably, despite its overfitting, the model still retains a robust accuracy of 40% under the L_2 attack—33% higher than that of the Standard model.




Fig. 7. Dynamic behavior of decision boundaries in the Rice2020Overfitting model during attack.

B. Defense Evaluation

The results in Tables 1 and 2 illustrate the impact of various input preprocessing defenses on both clean and adversarial accuracy across three model types. In the absence of any defense, all models experience a dramatic drop in adversarial accuracy, with the Standard model accuracy dropping to 0%. This highlights the inherent vulnerability of standard training to CW attacks, as also reported by XXX[3].

TABLE I
CLEAN ACCURACY (%) BY MODEL AND DEFENSE METHOD.
ABBREVIATIONS: STD (STANDARD), WONG (WONG2020FAST), RICE (RICE2020OVERFITTING).

	STD	WONG	RICE
No defense	96.0	92.0	86.0
Gaussian noise	93.4	84.1	86.1
Gray scale	92.1	78.7	80.7
JPEG compression	88.0	86.0	84.0
Randomized smoothing	32.0	88.0	86.0

TABLE II
ADVERSARIAL ACCURACY (%) BY MODEL AND DEFENSE METHOD.
ABBREVIATIONS: STD (STANDARD), WONG (WONG2020FAST), RICE (RICE2020OVERFITTING).

	STD	WONG	RICE
No defense	0.0	20.0	20.0
Gaussian noise	10.0	30.0	30.0
Gray scale	10.0	60.0	50.0
JPEG compression	68.0	56.0	62.0
Randomized smoothing	24.0	58.0	60.0

Introducing input transformations yields a moderate trade-off between clean and adversarial accuracy, consistent with findings by XXX [dziugaite2016study](#), [xu2018feature](#). While clean accuracy is slightly reduced for most defenses, adversarial accuracy improves substantially. For instance, JPEG compression reduces the clean accuracy of the Standard model from 96.0% to 88.0%, while boosting adversarial accuracy from 0.0% to 68.0%. This pattern holds across all three model architectures, indicating that JPEG compression is the most effective and consistent defense in this setting.

Randomized smoothing presents a unique behavior: it is highly effective for robustly trained models (WONG and RICE), but fails to offer significant protection for the Standard model, whose clean accuracy also drops sharply to 32.0%. This suggests that smoothing alone is not sufficient without robust model training.

Gaussian noise injection and grayscale transformations provide only modest improvements in adversarial accuracy, especially for the Standard model, which remains highly susceptible to attack. However, for WONG and RICE, these defenses yield intermediate adversarial accuracy, indicating some synergy with robust training.

Overall, among all tested defenses, **JPEG compression consistently achieves the best trade-off** between retaining high clean accuracy and providing substantial robust-




Fig. 8. Clean accuracy by different model and defense method.




Fig. 9. Adversarial accuracy by different model and defense methods.

ness to CW attacks. This superiority may be attributed to JPEG’s ability to remove high-frequency components and subtle pixel-level perturbations that adversarial attacks typically exploit XXX [dziugaite2016study](#), [guo2018countering](#). Consequently, JPEG compression emerges as a strong baseline for practical defenses, especially in scenarios where model retraining is not feasible.

V. SCENARIO APPLICATIONS

Adversarial attacks present significant risks across various real-world domains, underscoring the need for robust and practical defense mechanisms. Key application scenarios include:

- **Autonomous Driving:** Traffic sign recognition systems are highly vulnerable to adversarial attacks such as CW attacks. Deploying real-time defenses like JPEG compression can enhance the resilience of these systems, thereby promoting safety in autonomous vehicles.
- **Medical Imaging:** Adversarial perturbations in medical image classification (e.g., X-ray analysis) can lead to incorrect diagnoses. Robust input preprocessing methods are essential to safeguard clinical decision-making and ensure diagnostic reliability.
- **Federated Learning:** In collaborative learning settings, adversarially poisoned gradients can disrupt model integrity. Defense strategies, including dynamic aggregation algorithms and differential privacy, are pivotal for maintaining secure and trustworthy federated learning.

VI. CONCLUSION

We systematically evaluated input transformation-based defenses against CW adversarial attacks across three model types.

- **Attack Dynamics:** CW attacks devastate the Standard model (accuracy drops to 0%), while robust models (WONG and RICE) retain significantly higher adversarial accuracy.
- **Defense Effectiveness:** JPEG compression consistently offers the best trade-off, substantially boosting adversarial robustness with only moderate impact on clean accuracy. For instance, it raises adversarial accuracy from 0% to 68% in the Standard model.
- **Key Insights:** Input transformations, especially those filtering high-frequency perturbations, can significantly enhance robustness. However, some defenses (like randomized smoothing) are much less effective for non-robust models.

Overall, our results show that practical defenses like JPEG compression can provide strong protection against advanced adversarial attacks, especially when retraining models is not feasible.

CODE AVAILABILITY

The source code and reproducible scripts for all experiments in this work are publicly available at:
<https://github.com/bkty1122/COM6103>

REFERENCES

- [1] K. Eykholt, I. Evtimov, E. Fernandes, *et al.*, “Robust physical-world attacks on deep learning visual classification,” in *2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2018, pp. 1625–1634. DOI: 10.1109/CVPR.2018.00175.
- [2] S. G. Finlayson, H. W. Chung, I. S. Kohane, and A. L. Beam, “Adversarial Attacks Against Medical Deep Learning Systems,” *arXiv e-prints*, arXiv:1804.05296, arXiv:1804.05296, Apr. 2018. DOI: 10.48550/arXiv.1804.05296. arXiv: 1804.05296.
- [3] N. Carlini and D. Wagner, *Towards evaluating the robustness of neural networks*, 2017. arXiv:1608.04644 [cs.CR]. [Online]. Available: <https://arxiv.org/abs/1608.04644>.
- [4] I. Goodfellow, J. Shlens, and C. Szegedy, “Explaining and harnessing adversarial examples,” *arXiv 1412.6572*, Dec. 2014.
- [5] A. Madry, A. Makelov, L. Schmidt, D. Tsipras, and A. Vladu, *Towards deep learning models resistant to adversarial attacks*, 2019. arXiv: 1706.06083 [stat.ML]. [Online]. Available: <https://arxiv.org/abs/1706.06083>.